

A RANDOM WALK THROUGH RISK

Gambler's Ruin

Three betting scenarios — and the tactics that crack them

From a biased coin → closed-form formulas, clever code, a little linear algebra

3 scenarios • 1 underlying random walk



Some common upcoming terms

- **States:** The distinct conditions or situations the system can be in (e.g., money at an instant).
- **Transitions:** The movement of the system from one state to another, or from a state back to itself (self-loop).
- **Transition Probability:** The mathematical likelihood (between 0 and 1) that the system will move from a given current state to a specific future state.
- **Markov Chain:** A mathematical system that models a sequence of events where the probability of the next step depends solely on the **current state**, not on how the process arrived there. This defining feature is known as the “**Markov property**” or “**memorylessness**”.
- **Stationary Distribution:** A probability distribution over the states of a system that remains unchanged as the system evolves over time. (**Note:** If a system begins in this distribution, its state probabilities will remain constant indefinitely.)
- **Transition Matrix:** A square matrix where rows and columns represent states. Each element A_{ij} depicts the probability to go from state i to state j .

Why should an ML team care about a gambler?

Every system below is the **same mathematical object** in a different costume: a process that hops between states, where the next step depends only on where you are *now*.

Reinforcement Learning

An agent choosing actions *is* a gambler choosing bets. Value iteration solves it.

Google PageRank

A surfer clicking links at random; page rank = the stationary distribution.

Reliability & Networks

Time-to-failure, churn, how fast something spreads = a first-hitting time.

Quant Finance

Stock prices, ruin probability, risk of bankruptcy.

Big idea: learn it once on a toy gambling problem, and a good part of modern ML comes for free.

One coin, one number line

Every scenario is the **same picture**: a token on a number line that steps **right (prob. p)** or **left (prob. $q = 1 - p$)** each round. All that changes is the rule at the edges — and the question we ask.



Why it's solvable: the Markov property

Where you go next depends *only* on where you stand now — never on how you got there. Memoryless.

The one trick we reuse everywhere

Condition on the very first step:

$$f(\text{here}) = p f(\text{step right}) + q f(\text{step left}) (+ \text{cost})$$

Three scenarios

Same random walk, three disguises. We'll state each puzzle, then crack it — watching *which* tactic the problem rewards.

1 The Classic Gambler

Fixed \$1 bets. Her odds, her ruin, her time at the table.

2 The Bold Gambler

All-in betting

3 The Stock Market

A drifting price: long-run behaviour & first hits.

One idea, three costumes — and a different tool for each.

01

SCENARIO 1

The Classic Gambler

A biased coin, a target fortune, and one wall she never wants to hit.

The setup

She walks in with $\$k$. Every round she stakes $\$1$ on a biased coin: heads (prob. p) and she's up a dollar; tails (prob. q) and she's down one. She stands up the instant she hits her **target $\$N$** — or busts out at **$\0** .

Three things we want to know


1. What's her chance of reaching $\$N$ before going broke?
2. If she ditches the target and just plays until ruin — can she get *infinitely* rich?
3. How long will she sit at the table before it's over?

She starts somewhere between the two walls — and the walls decide her fate.

Cracking it — her odds of winning


1. **Condition on the first bet.** From $\$k$ she lands on $\$k+1$ (prob. p) or $\$k-1$ (prob. q):

$$T(k) = pT(k+1) + qT(k-1).$$



It all rides on q/p . A closed form is the dream:
instant, exact, no computer needed.

Cracking it — her odds of winning



1. **Condition on the first bet.** From $\$k$ she lands on $\$k+1$ (prob. p) or $\$k-1$ (prob. q):


$$T(k) = pT(k+1) + qT(k-1).$$

2. **Solve it.** Characteristic eq: $px^2 - x + q = 0$, roots 1 and q/p . General solution:

$$T(k) = A(1)^k + B(q/p)^k$$

It all rides on q/p . A closed form is the dream:
instant, exact, no computer needed.

Cracking it — her odds of winning



1. **Condition on the first bet.** From $\$k$ she lands on $\$k+1$ (prob. p) or $\$k-1$ (prob. q):

$$T(k) = pT(k+1) + qT(k-1).$$

2. **Solve it.** Characteristic eq: $px^2 - x + q = 0$, roots 1 and q/p . General solution:

$$T(k) = A(1)^k + B(q/p)^k$$


3. **Fit the ends:** Use $T(0) = 0$ and $T(N) = 1$ to solve for A and B :

$$p \neq q : T(k) = \frac{1 - (q/p)^k}{1 - (q/p)^N}; \quad p = q : T(k) = \frac{k}{N}.$$


It all rides on q/p . A closed form is the dream:
instant, exact, no computer needed.

3 gamblers 3 coins


Each line is one simulated life: start at wealth 25, walls at 0 and 50. Same rules, wildly different fates.



slight disadvantage drags everyone to ruin



a coin-flip of destinies



a slight edge, most paths climb

Cracking it — forever, and for how long?

Can she get infinitely rich?

Let $N \rightarrow \infty$.

Fair or against her ($p \leq \frac{1}{2}$): ruin is **certain**, $T \rightarrow 0$.

Only with a true edge ($p > \frac{1}{2}$): $T \rightarrow 1 - (q/p)^k > 0$.

This is “risk of ruin” — why the house, on its sliver of edge, takes everyone eventually.

How long at the table?

Same trick, +1 for the round just played:

$$E_i = p E_{i+1} + q E_{i-1} + 1, \quad E_0 = E_N = 0.$$

Solving this gives two distinct cases: **If $p \neq q$:**

$$E_i = \frac{i}{q-p} - \frac{N}{q-p} \frac{1 - (q/p)^i}{1 - (q/p)^N}$$

If $p = q$ (Fair Game):

$$E_i = i(N - i)$$

The fair game lasts longest right in the middle.

Where this bites in real life: position sizing, insurance reserves, bankroll management.

02

SCENARIO 2

The Bold Gambler

Tired of nickel-and-dime bets, she starts going all in.

The setup

New attitude, new rule. With **less than half** the target ($k < \frac{N}{2}$) she shoves **all** her chips — double up (prob. p) or bust (prob. q). With **at least half** ($k \geq \frac{N}{2}$) she bets exactly enough to finish, $N - k$: hit the target (prob. p), or slip back to $2k - N$ (prob. q) and keep going.

Probability of winning: $P(x)$

$$P(x) = \begin{cases} 0 & x = 0 \\ pP(2x) & 0 < x < \frac{N}{2} \\ qP(2x - N) + p & \frac{N}{2} \leq x < N \\ 1 & x = N \end{cases}$$

Expected duration: $\mathbb{E}(x)$

$$\mathbb{E}(x) = \begin{cases} 0 & x = 0 \\ 1 + p\mathbb{E}(2x) & 0 < x < \frac{N}{2} \\ 1 + q\mathbb{E}(2x - N) & \frac{N}{2} \leq x < N \\ 0 & x = N \end{cases}$$

Both the odds of winning and expected time follow similar piecewise recursions.

Cracking it — Probability via Binary

1. **Spot the binary.** Represent the ratio $\frac{k}{N}$ in binary form: $0.n_1n_2n_3 \dots$

The takeaway

Instead of a complicated formula, the recursion simplifies into a straightforward traversal of the binary string representing the wealth fraction.

Cracking it — Probability via Binary

1. **Spot the binary.** Represent the ratio $\frac{k}{N}$ in binary form: $0.n_1n_2n_3 \dots$
2. **Shift operations.** Doubling $2x$ is a *left shift* of the digits. Losing from the top ($2x - N$) shifts *and drops* the leading 1. So the game just reads the binary string bit by bit.

The takeaway

Instead of a complicated formula, the recursion simplifies into a straightforward traversal of the binary string representing the wealth fraction.

Cracking it — Probability via Binary

1. **Spot the binary.** Represent the ratio $\frac{k}{N}$ in binary form: $0.n_1n_2n_3 \dots$
2. **Shift operations.** Doubling $2x$ is a *left shift* of the digits. Losing from the top ($2x - N$) shifts *and drops* the leading 1. So the game just reads the binary string bit by bit.
3. **Walk the bits.** To calculate $P(x)$, traverse the binary representation. For the i^{th} bit:

The takeaway

Instead of a complicated formula, the recursion simplifies into a straightforward traversal of the binary string representing the wealth fraction.

Cracking it — Probability via Binary

1. **Spot the binary.** Represent the ratio $\frac{k}{N}$ in binary form: $0.n_1n_2n_3 \dots$
2. **Shift operations.** Doubling $2x$ is a *left shift* of the digits. Losing from the top ($2x - N$) shifts *and drops* the leading 1. So the game just reads the binary string bit by bit.
3. **Walk the bits.** To calculate $P(x)$, traverse the binary representation. For the i^{th} bit:
 - If **bit is 0**: Multiply the running probability by p (equivalent to $pP(2x)$).

The takeaway

Instead of a complicated formula, the recursion simplifies into a straightforward traversal of the binary string representing the wealth fraction.

Cracking it — Probability via Binary

1. **Spot the binary.** Represent the ratio $\frac{k}{N}$ in binary form: $0.n_1n_2n_3 \dots$
2. **Shift operations.** Doubling $2x$ is a *left shift* of the digits. Losing from the top ($2x - N$) shifts *and drops* the leading 1. So the game just reads the binary string bit by bit.
3. **Walk the bits.** To calculate $P(x)$, traverse the binary representation. For the i^{th} bit:
 - If **bit is 0**: Multiply the running probability by p (equivalent to $pP(2x)$).
 - If **bit is 1**: Multiply by q and add p (equivalent to $qP(2x - N) + p$).

The takeaway

Instead of a complicated formula, the recursion simplifies into a straightforward traversal of the binary string representing the wealth fraction.

Cracking it — Expected Duration via Binary

1. **Apply the same logic.** The expected duration $\mathbb{E}(x)$ shares the exact same state transitions ($2x$ and $2x - N$). We can parse the exact same binary string to solve it.

The takeaway

The entire strategy condenses into an elegant $O(\log N)$ bit-manipulation loop. A hard problem dissolves when viewed through the right data representation.

Cracking it — Expected Duration via Binary

1. **Apply the same logic.** The expected duration $\mathbb{E}(x)$ shares the exact same state transitions ($2x$ and $2x - N$). We can parse the exact same binary string to solve it.
2. **Walk the bits.** Traverse the binary representation of $\frac{k}{N}$. For the i^{th} bit:

The takeaway

The entire strategy condenses into an elegant $O(\log N)$ bit-manipulation loop. A hard problem dissolves when viewed through the right data representation.

Cracking it — Expected Duration via Binary

1. **Apply the same logic.** The expected duration $\mathbb{E}(x)$ shares the exact same state transitions ($2x$ and $2x - N$). We can parse the exact same binary string to solve it.
2. **Walk the bits.** Traverse the binary representation of $\frac{k}{N}$. For the i^{th} bit:
 - If **bit is 0**: Multiply the expectation by p and add 1 (equivalent to $1 + p\mathbb{E}(2x)$).

The takeaway

The entire strategy condenses into an elegant $O(\log N)$ bit-manipulation loop. A hard problem dissolves when viewed through the right data representation.

Cracking it — Expected Duration via Binary

1. **Apply the same logic.** The expected duration $\mathbb{E}(x)$ shares the exact same state transitions ($2x$ and $2x - N$). We can parse the exact same binary string to solve it.
2. **Walk the bits.** Traverse the binary representation of $\frac{k}{N}$. For the i^{th} bit:
 - If **bit is 0**: Multiply the expectation by p and add 1 (equivalent to $1 + p\mathbb{E}(2x)$).
 - If **bit is 1**: Multiply the expectation by q and add 1 (equivalent to $1 + q\mathbb{E}(2x - N)$).

The takeaway

The entire strategy condenses into an elegant $O(\log N)$ bit-manipulation loop. A hard problem dissolves when viewed through the right data representation.

03


SCENARIO 3

The Stock Market

Same random walk, new costume — a drifting price.

The setup

Step away from the casino and watch a stock. Each day the price ticks **up one** (prob. p_k), **down one** (prob. q_k), or **holds** (prob. r_k) — rates that can depend on the current price k . It lives on $\{0, 1, \dots, N\}$ and can't escape the ends.



Two questions

1. Long-run: how much time does the price spend at each level, and what's the *average* price?
2. From price a , how long until it first touches price b ?

Cracking it — the long-run picture

1. Long-run proportions π are unchanged by one more step:
 $\pi = \pi P$, where:

$$P = \begin{bmatrix} r_0 & p_0 & 0 & \dots & 0 \\ q_1 & r_1 & p_1 & \dots & 0 \\ 0 & q_2 & r_2 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & q_N & r_N \end{bmatrix}$$



Where it spends its time — and the start point is forgotten.

Closed-form ratios, then one normalisation — a few lines of cumulative products.

Cracking it — the long-run picture

1. Long-run proportions π are unchanged by one more step:
 $\pi = \pi P$, where:

$$P = \begin{bmatrix} r_0 & p_0 & 0 & \dots & 0 \\ q_1 & r_1 & p_1 & \dots & 0 \\ 0 & q_2 & r_2 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & q_N & r_N \end{bmatrix}$$

2. **Balance neighbours:** up-flow = down-flow,

$$\pi_{i-1}p_{i-1} = \pi_i q_i \Rightarrow \pi_i = \pi_{i-1} \frac{p_{i-1}}{q_i}.$$

Closed-form ratios, then one normalisation — a few lines of cumulative products.



Where it spends its time — and the start point is forgotten.

Cracking it — the long-run picture

1. Long-run proportions π are unchanged by one more step:
 $\pi = \pi P$, where:


$$P = \begin{bmatrix} r_0 & p_0 & 0 & \dots & 0 \\ q_1 & r_1 & p_1 & \dots & 0 \\ 0 & q_2 & r_2 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & q_N & r_N \end{bmatrix}$$

2. **Balance neighbours:** up-flow = down-flow,

$$\pi_{i-1}p_{i-1} = \pi_i q_i \Rightarrow \pi_i = \pi_{i-1} \frac{p_{i-1}}{q_i}.$$

3. Chain the ratios from π_0 ; pin π_0 with $\sum_i \pi_i = 1$.

Closed-form ratios, then one normalisation — a few lines of cumulative products.



Where it spends its time — and the start point is forgotten.

Cracking it — the long-run picture

1. Long-run proportions π are unchanged by one more step:
 $\pi = \pi P$, where:

$$P = \begin{bmatrix} r_0 & p_0 & 0 & \dots & 0 \\ q_1 & r_1 & p_1 & \dots & 0 \\ 0 & q_2 & r_2 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & q_N & r_N \end{bmatrix}$$

2. **Balance neighbours:** up-flow = down-flow,

$$\pi_{i-1} p_{i-1} = \pi_i q_i \Rightarrow \pi_i = \pi_{i-1} \frac{p_{i-1}}{q_i}.$$

3. Chain the ratios from π_0 ; pin π_0 with $\sum_i \pi_i = 1$.

4. Average price = $\sum_i i \pi_i$.

Closed-form ratios, then one normalisation — a few lines of cumulative products.



Where it spends its time — and the start point is forgotten.

Cracking it — time to hit a target price

1. Let $\mathbb{E}_x$ be the expected days from x to b . Condition on the first move:

$$\mathbb{E}_x = p_x \mathbb{E}_{x+1} + r_x \mathbb{E}_x + q_x \mathbb{E}_{x-1} + 1, \quad \mathbb{E}_b = 0.$$

*Real-world: that stationary π is **PageRank**; a first-hitting time is expected **time-to-failure**, **time-to-churn**, or how long something takes to spread.*

Cracking it — time to hit a target price

1. Let $\mathbb{E}_x$ be the expected days from x to b . Condition on the first move:

$$\mathbb{E}_x = p_x \mathbb{E}_{x+1} + r_x \mathbb{E}_x + q_x \mathbb{E}_{x-1} + 1, \quad \mathbb{E}_b = 0.$$

2. One equation per state. Stack them into a linear system $A\mathbf{E} = \mathbf{1}$:

$$\begin{bmatrix} 1 - r_0 & -p_0 & 0 & \dots & 0 \\ -q_1 & 1 - r_1 & -p_1 & \dots & 0 \\ 0 & -q_2 & 1 - r_2 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 1 & \dots \\ \vdots & \vdots & \vdots & \vdots & \ddots \\ 0 & 0 & \dots & -q_N & 1 - r_N \end{bmatrix} \begin{bmatrix} \mathbb{E}_0 \\ \mathbb{E}_1 \\ \mathbb{E}_2 \\ \vdots \\ \mathbb{E}_b \\ \vdots \\ \mathbb{E}_N \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \\ \vdots \\ 0 \\ \vdots \\ 1 \end{bmatrix}$$

Real-world: that stationary π is **PageRank**; a first-hitting time is expected **time-to-failure**, **time-to-churn**, or how long something takes to spread.

Cracking it — time to hit a target price

1. Let $\mathbb{E}_x$ be the expected days from x to b . Condition on the first move:

$$\mathbb{E}_x = p_x \mathbb{E}_{x+1} + r_x \mathbb{E}_x + q_x \mathbb{E}_{x-1} + 1, \quad \mathbb{E}_b = 0.$$

2. One equation per state. Stack them into a linear system $A\mathbf{E} = \mathbf{1}$:

$$\begin{bmatrix} 1 - r_0 & -p_0 & 0 & \dots & 0 \\ -q_1 & 1 - r_1 & -p_1 & \dots & 0 \\ 0 & -q_2 & 1 - r_2 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 1 & \dots \\ \vdots & \vdots & \vdots & \vdots & \ddots \\ 0 & 0 & \dots & -q_N & 1 - r_N \end{bmatrix} \begin{bmatrix} \mathbb{E}_0 \\ \mathbb{E}_1 \\ \mathbb{E}_2 \\ \vdots \\ \mathbb{E}_b \\ \vdots \\ \mathbb{E}_N \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \\ \vdots \\ 0 \\ \vdots \\ 1 \end{bmatrix}$$

3. Notice row b is pinned entirely: $1 \cdot \mathbb{E}_b = 0$. Since the rates vary, there's no clean formula — hand the matrix to a solver.

*Real-world: that stationary π is **PageRank**; a first-hitting time is expected **time-to-failure**, **time-to-churn**, or how long something takes to spread.*


Three ways to crack a recurrence

Notice we never used one hammer for everything. The puzzle tells you which tool is cheapest.

- 1 **Closed form** (*Scenario 1*) — constant coefficients → characteristic equation → a formula. Instant, exact, hands you insight (that q/p ratio).
- 2 **Walk the recursion** (*Scenario 2*) — formula's ugly, but the states are structured → a tiny loop.
- 3 **Linear system** (*Scenario 3*) — coefficients vary → "it's just $A\mathbf{E} = \mathbf{b}$," hand it to a solver.

Always in your back pocket

Simulate. Stuck? Run the random walk thousands of times and average. Approximate — but it never fails to apply.




Reach for the cheapest tool that works — and know when to stop hunting for a formula.

Beyond the casino

Strip off the story and this same random walk is everywhere.

- **Risk of ruin** — position sizing, insurance reserves and bankroll math are exactly Scenario 1's ruin probability.
- **Learning the best bet** — let the gambler *choose* her stake and you've built reinforcement learning; the optimal play is bold, with a famously fractal value curve (right).
- **Reliability & ops** — a first-hitting time is expected *time-to-failure*, time-to-churn, or how fast a rumour, infection, or queue backlog spreads.
- **PageRank** — the stationary distribution of a random web-surfer; the next slide unpacks it.




The same machinery quietly powers a lot of modern ML.

PageRank — the random surfer

Swap the number line for a **graph of web pages**. A surfer stands on a page and, each step, clicks a *random* outgoing link. Run this forever: the fraction of time spent on each page *is* its rank. Same stationary-distribution math as the drifting stock — just on a messier state space.

1. **Same balance condition.** The rank vector is fixed under one more click: $\pi = \pi M$, where M_{ij} is the chance of hopping $i \rightarrow j$.
2. **Vote by linking.** A page is important if important pages link to it — importance flows along edges, split evenly across a page's out-links.
3. **The teleport fix.** A surfer can get stuck on dead-end or cyclic pages. With prob. $1-d$ (≈ 0.15) she **teleports** to a random page. That damping makes the chain *irreducible & aperiodic* — so a *unique* π exists and is reached from any start.
4. **Solve by iterating.** Start uniform, apply M repeatedly — power iteration converges fast on a web-scale, billion-node graph.



A is biggest: three pages point at it. Node size $\propto$ rank.

Why it mattered

Treating links as votes — and ranking by the random walk's stationary distribution — is the idea that launched Google. The same eigenvector trick now drives citation ranking, recommendation graphs, and fraud-ring detection.

Let's talk

This deck is built to **pause at each scenario** — predict the next step before the click. A few things worth arguing about:

- Even at fair odds, distance-to-target crushes you — why does that feel so counterintuitive?
- When is hunting for a closed form a waste of time versus just coding the recurrence?
- Why do changing transition probabilities (p_k, q_k) force us into linear algebra instead of closed forms?
- Bold play beating cautious play under bad odds — does that match your gut on risk?
- “Where does this spend most of its time?” — where else is that the useful question?
- What quantity in your world is secretly a *first hitting time*?

THANK YOU

**One coin.
Three scenarios.
A handful of tactics.**

Questions?