

A RANDOM WALK THROUGH RISK

Gambler's Ruin

Why casinos always win, and other stories from one humble coin

From a biased coin → closed-form formulas, clever code, a little linear algebra

3 scenarios • 3 audience polls • 1 underlying random walk



Let's open with a bet

You're in a casino with **\$100**. You bet **\$1 on red** every spin of the roulette wheel, and you walk away the moment you either **double up to \$200** — or **lose everything**. Each spin, you win with probability $18/38 \approx 47.4\%$. Only a whisker below a fair coin.

What's your chance of walking out with \$200? *(Any logical guesses?)*

A ~47%

B ~23%

C ~3%

D <0.01%

Let's open with a bet

You're in a casino with **\$100**. You bet **\$1 on red** every spin of the roulette wheel, and you walk away the moment you either **double up to \$200** — or **lose everything**. Each spin, you win with probability $18/38 \approx 47.4\%$. Only a whisker below a fair coin.

What's your chance of walking out with \$200? *(Any logical guesses?)*

A ~47%

B ~23%

C ~3%

D <0.01%

The answer is D — about 1 in 37,000

A 47.4% game *per spin* collapses to a **0.003%** chance *overall*. The casino's tiny 2.6-point edge, compounded across your "safe" little \$1 bets, is a near-certain death sentence. By the end of this session you'll compute this in one line — **and know the one change that multiplies those odds 18,000-fold.**

This one theorem moves serious money

Casinos

\$650 billion/year global revenue

House edge: just **2–5%**

Gambler's ruin turns that into near-certain profit

Insurance

Insurer $\sim$ Gambler (premiums in, claims out)

EU Solvency II mandate:

P(ruin within 1 year) < 0.5%

Retail Trading

SEBI study (FY22–24):

93% of Indian F&O traders lost money

Aggregate losses: **Rs. 1.8 lakh crore**

Google PageRank


PageRank $\sim$ Stationary distribution of a random walk

Seeded a company now worth **\$3+ trillion**

One framework, four industries — let's build the mathematics.

What is Gambler's Ruin?

- **The picture:** A player walks in with a starting amount, making the same bet over and over until they hit a target or go broke. Think of a token bouncing randomly between two trapdoors on a number line; while we cannot predict a single token's actual path, we can perfectly map the system's theoretical probabilities and expected lifespans.
- **Why you care:** dress the token up differently and it becomes an insurer's capital, a stock price, or a web surfer. Crack one, and you've cracked them all.



The only jargon you'll need today

- **States:** The distinct conditions or situations the system can be in (e.g., money at an instant).
- **Transitions:** The movement of the system from one state to another, or from a state back to itself (self-loop).
- **Transition Probability:** The mathematical likelihood (between 0 and 1) that the system will move from a given current state to a specific future state.
- **Markov Chain:** A mathematical system that models a sequence of events where the probability of the next step depends solely on the **current state**, not on how the process arrived there. This defining feature is known as the “**Markov property**” or “**memorylessness**”.
- **Stationary Distribution:** A probability distribution over the states of a system that remains unchanged as the system evolves over time. (**Note:** If a system begins in this distribution, its state probabilities will remain constant indefinitely.)
- **Transition Matrix:** A square matrix where rows and columns represent states. Each element A_{ij} depicts the probability to go from state i to state j .

THE COMMON THREAD

One coin, one number line

Every scenario is the **same picture**: a token on a number line that steps **right (prob. p)** or **left (prob. $q = 1 - p$)** each round. All that changes is the rule at the edges — and the question we ask.



Why it's solvable: the Markov property

Where you go next depends *only* on where you stand now — never on how you got there. Memoryless.

The one trick we reuse everywhere

Condition on every step:

$$T(k) = p(T(k+1) + cr) + q(T(k-1) + cl)$$

Three scenarios

Same random walk, three disguises. We'll state each puzzle, then crack it — watching *which* tactic the problem rewards.

1 The Classic Gambler

Fixed \$1 bets. Her odds, her ruin, her time at the table.

2 The Bold Gambler

All-in, double-or-die betting. Reckless? The math will disagree.

3 The Stock Market

A drifting price: long-run behaviour & first hits.

One idea, three costumes — and a different tool for each. Two polls still ahead: keep your hands warm.

01

SCENARIO 1

The Classic Gambler

A biased coin, a target fortune, and one wall she never wants to hit.

The setup

She walks in with $\$k$. Every round she stakes \$1 on a biased coin: heads (prob. p) and she's up a dollar; tails (prob. q) and she's down one. She stands up the instant she hits her **target $\$N$** — or is broke at **$\0** .

Three things we want to know

1. What's her chance of reaching $\$N$ before going broke?
2. If she ditches the target and just plays until ruin — can she get *infinitely* rich?
3. How long will she sit at the table before it's over?

She starts somewhere between the two walls — and the walls decide her fate.

Cracking it — her odds of winning

1. **Condition on the first bet.** From $\$k$ she lands on $\$k+1$ (prob. p) or $\$k-1$ (prob. q):

$$T(k) = pT(k+1) + qT(k-1).$$

2. **Rearrange.** Since $p + q = 1$, rewrite as:

$$p[T(k+1) - T(k)] = q[T(k) - T(k-1)].$$


Let $D(k) = T(k+1) - T(k)$. Then $D(k) = \frac{q}{p} D(k-1)$.

Each step-difference is $\frac{q}{p}$ times the previous one — **a geometric sequence!**

$$D(k) = \left(\frac{q}{p}\right)^k D(0).$$

3. **Telescope & fit boundaries.** $T(k) = T(0) + \sum_{j=0}^{k-1} D(j)$ is a geometric sum. Plug in $T(0) = 0$, $T(N) = 1$:


$$p \neq q : T(k) = \frac{1 - (q/p)^k}{1 - (q/p)^N}; \quad p = q : T(k) = \frac{k}{N}.$$




It all rides on q/p . A closed form is the dream:
instant, exact, no computer needed.

3 gamblers 3 coins


Each line is one simulated life: start at wealth 25, walls at 0 and 50. Same rules, wildly different fates.



slight disadvantage drags everyone to ruin



a coin-flip of destinies



a slight edge, most paths climb

The fair-coin trap

Now tweak the story: the coin is **perfectly fair** ($p = \frac{1}{2}$), and we **remove the target** — she just keeps playing. Forever, if she can. What happens in the long run?

A Drifts around, breaks even

B 50–50: rich or ruined

C Ruin is *certain*

The fair-coin trap

Now tweak the story: the coin is **perfectly fair** ($p = \frac{1}{2}$), and we **remove the target** — she just keeps playing. Forever, if she can. What happens in the long run?

- A Drifts around, breaks even B 50–50: rich or ruined C Ruin is *certain*

C. With a fair coin and no exit, ruin has probability 1.

The walls aren't symmetric: **\$0 is a trapdoor** she can fall through, but infinity is not a wall she can ever touch. And the cruellest twist: the *expected* time until she goes broke is **infinite**. Certain doom — on no schedule. Proof on the next slide.

Cracking it — forever, and for how long?

Can she get infinitely rich?

Let $N \rightarrow \infty$.

Fair or against her ($p \leq \frac{1}{2}$): ruin is **certain**, $T \rightarrow 0$.

Only with a true edge ($p > \frac{1}{2}$): $T \rightarrow 1 - (q/p)^k > 0$.

This is “risk of ruin” — why the house, on its sliver of edge, takes everyone eventually.

Why the +1? Each step costs one round:

$$E_i = p \underbrace{[E_{i+1} + 1]}_{\text{win}} + q \underbrace{[E_{i-1} + 1]}_{\text{lose}}$$

Expand and use $p + q = 1$:

$$E_i = p E_{i+1} + q E_{i-1} + \underbrace{(p + q)}_{=1}$$

How long at the table?

Same trick, +1 for the round just played:

$$E_i = p E_{i+1} + q E_{i-1} + 1, \quad E_0 = E_N = 0.$$

Solving this gives two distinct cases: **If $p \neq q$:**

$$E_i = \frac{i}{q - p} - \frac{N}{q - p} \frac{1 - (q/p)^i}{1 - (q/p)^N}$$

If $p = q$ (Fair Game):

$$E_i = i(N - i)$$

The fair game lasts longest right in the middle.

Where this bites in real life: position sizing, insurance reserves, bankroll management.

02

SCENARIO 2

The Bold Gambler

Tired of nickel-and-dime bets, she starts going all in.

Pick your strategy

Back to our opening casino: **47.4%** odds per round, you hold **\$100**, and only **\$200** gets you home — anything less is worthless to you. Two ways to play:

A Timid

Bet \$1 at a time. Stay in the game, grind it out, let the law of averages be your friend.

B Bold

Shove everything on one spin. Double or die. Repeat only if you must.

Who plays timid? Who goes bold?

Pick your strategy

Back to our opening casino: **47.4%** odds per round, you hold **\$100**, and only **\$200** gets you home — anything less is worthless to you. Two ways to play:

A Timid

Bet \$1 at a time. Stay in the game, grind it out, let the law of averages be your friend.

B Bold

Shove everything on one spin. Double or die. Repeat only if you must.

Who plays timid? Who goes bold?

Hold that vote.

You already know the timid number from the opening poll: **1 in 37,000**. Can recklessness possibly beat it? Let's define bold play properly, crack it with a trick you'll genuinely enjoy — then settle the score.

The setup

New attitude, new rule. With **less than half** the target ($k < \frac{N}{2}$) she shoves **all** her chips — double up (prob. p) or broke (prob. q). With **at least half** ($k \geq \frac{N}{2}$) she bets exactly enough to finish, $N - k$: hit the target (prob. p), or slip back to $2k - N$ (prob. q) and keep going.

Probability of winning: $P(x)$

$$P(x) = \begin{cases} 0 & x = 0 \\ pP(2x) & 0 < x < \frac{N}{2} \\ qP(2x - N) + p & \frac{N}{2} \leq x < N \\ 1 & x = N \end{cases}$$

Expected duration: $\mathbb{E}(x)$

$$\mathbb{E}(x) = \begin{cases} 0 & x = 0 \\ 1 + p\mathbb{E}(2x) & 0 < x < \frac{N}{2} \\ 1 + q\mathbb{E}(2x - N) & \frac{N}{2} \leq x < N \\ 0 & x = N \end{cases}$$

Both the odds of winning and expected time follow similar piecewise recursions.

Cracking it — Probability via Binary

1. **Why binary at all?** Write the wealth fraction $\frac{x}{N}$ in binary: $0.b_1b_2b_3\dots$. The first bit b_1 answers exactly the question the game asks first — “is x in the lower half ($b_1=0$) or upper half ($b_1=1$) of $[0, N]$?”
2. **The two moves are bit-shifts in disguise.**
 - Lower half ($b_1=0, x < \frac{N}{2}$): she doubles, $x \rightarrow 2x$. The fraction $\frac{2x}{N}$ is just $\frac{x}{N}$ *shifted one bit left*; the leading 0 falls off, b_2 becomes the new lead bit.
 - Upper half ($b_1=1, x \geq \frac{N}{2}$): on a loss she drops to $2x - N$. The fraction $\frac{2x-N}{N} = 2 \cdot \frac{x}{N} - 1$ — a left shift, but now the leading 1 is the one that falls off. Either way, **one bit gets consumed per round**.
3. **So the recursion just reads the binary string.** Walk $b_1b_2b_3\dots$ left-to-right, maintaining a running probability P . Each bit fires one of the two recurrence cases:
 - bit is 0: $P \leftarrow pP$ (matches $P(x) = pP(2x)$)
 - bit is 1: $P \leftarrow qP + p$ (matches $P(x) = qP(2x-N) + p$)

The takeaway

What looked like a tangled piecewise recursion is actually a deterministic tour of the bits of x/N . The right *representation* of the state space turns it into an $O(\log N)$ loop — no closed form required.

Cracking it — Expected Duration via Binary

1. **Same state machine, different ledger.** The moves are unchanged: lower half doubles ($x \rightarrow 2x$, a left shift dropping a 0), upper half either ends or drops to $2x - N$ (a left shift dropping a 1). So the binary digits of $\frac{x}{N}$ *still* drive the game — only what we accumulate along the way changes.
2. **Why each step adds $+1$.** We're now counting *rounds*, not winning probability. Every bit consumed corresponds to one bet that actually got played, so we add 1 each time. The multiplier in front of $\mathbb{E}$ is the chance the game *survives* that round (and so still has rounds left to count).
 - bit is 0 (lower half, she doubled): $\mathbb{E} \leftarrow 1 + p\mathbb{E}$ (game continues only on a win, prob. p)
 - bit is 1 (upper half, she played for the win): $\mathbb{E} \leftarrow 1 + q\mathbb{E}$ (game continues only on a loss, prob. q)
3. **When does it end?** The bit-walk terminates when we run out of bits — i.e. when $\frac{x}{N}$ has a finite binary expansion — exactly when the game has fully resolved.

The takeaway

Two completely different quantities — a probability and an expectation — fall to the *same* bit-walk, because both inherit the same state machine. When the recursion's state space matches a familiar data representation, the representation itself is the algorithm.

The scoreboard: bold vs. timid

Same casino, same \$100 in hand, same goal of \$200 — win probability $18/38 \approx 47.4\%$ per round.

Timid play (\$1 bets)

0.003%

About 1 in 37,000. The house edge gets ~ 200 swings at you
— and it uses every single one.

Bold play (all-in)

47.4%

One spin, one judgement. Luck only gets a single chance to
hurt you.

Bold play is $\approx 18,000\times$ more likely to succeed

Not a fluke — a theorem (Dubins & Savage, 1965): **when the odds are against you, bold play is optimal** — no strategy does better. Flip side: when the edge is *yours* ($p > \frac{1}{2}$), patient grinding wins — which is exactly how the *casino* plays *you*. **Moral:** underdogs should minimise how often luck gets to judge them; favourites should maximise it.

03


SCENARIO 3

The Stock Market

Same random walk, new costume — a drifting price.

The setup

Step away from the casino and watch a stock. Each day the price ticks **up one** (prob. p_k), **down one** (prob. q_k), or **holds** (prob. r_k) — rates that can depend on the current price k . It lives on $\{0, 1, \dots, N\}$ and can't escape the ends.



Two questions

1. Long-run: how much time does the price spend at each level, and what's the *average* price?
2. From price a , how long until it first touches price b ?

Cracking it — the long-run picture

1. Long-run proportions form a **row vector**

$$\pi = [\pi_0, \pi_1, \pi_2, \dots, \pi_N],$$

with $\pi_i \geq 0$ and $\sum_i \pi_i = 1$. It's the distribution that's unchanged by one more step: $\pi = \pi P$, where


$$P = \begin{bmatrix} r_0 & p_0 & 0 & \dots & 0 \\ q_1 & r_1 & p_1 & \dots & 0 \\ 0 & q_2 & r_2 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & q_N & r_N \end{bmatrix}$$

2. **Balance neighbours:** up-flow = down-flow,

$$\pi_{i-1}p_{i-1} = \pi_i q_i \Rightarrow \pi_i = \pi_{i-1} \frac{p_{i-1}}{q_i}.$$

3. Chain the ratios from π_0 ; pin π_0 with $\sum_i \pi_i = 1$.
4. Average price = $\sum_i i \pi_i$.

Closed-form ratios, then one normalisation.



Where it spends its time — and the start point is forgotten.

Cracking it — time to hit a target price

1. Let $\mathbb{E}_x$ be the expected days from x to b . Condition on the first move:

$$\mathbb{E}_x = p_x \mathbb{E}_{x+1} + r_x \mathbb{E}_x + q_x \mathbb{E}_{x-1} + 1, \quad \mathbb{E}_b = 0.$$

2. One equation per state. Stack them into a linear system $A\mathbf{E} = \mathbf{1}$:

$$\begin{bmatrix} 1 - r_0 & -p_0 & 0 & \dots & 0 \\ -q_1 & 1 - r_1 & -p_1 & \dots & 0 \\ 0 & -q_2 & 1 - r_2 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 1 & \dots \\ \vdots & \vdots & \vdots & \vdots & \ddots \\ 0 & 0 & \dots & -q_N & 1 - r_N \end{bmatrix} \begin{bmatrix} \mathbb{E}_0 \\ \mathbb{E}_1 \\ \mathbb{E}_2 \\ \vdots \\ \mathbb{E}_b \\ \vdots \\ \mathbb{E}_N \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \\ \vdots \\ 0 \\ \vdots \\ 1 \end{bmatrix}$$


3. Notice row b is pinned entirely: $1 \cdot \mathbb{E}_b = 0$. Since the rates vary, there's no clean formula — hand the matrix to a solver.

*Real-world: that stationary π is **PageRank**; a first-hitting time is expected **time-to-failure**, **time-to-churn**, or how long something takes to spread.*

PageRank — the random surfer

Swap the number line for a **graph of web pages**. A surfer stands on a page and, each step, clicks a *random* outgoing link. Run this forever: the fraction of time spent on each page *is* its rank. Same stationary-distribution math as the drifting stock — just on a messier state space.

1. **Same balance condition.** The rank vector is fixed under one more click: $\pi = \pi M$, where M_{ij} is the chance of hopping $i \rightarrow j$.
2. **Vote by linking.** A page is important if important pages link *to* it — importance flows along edges, split evenly across a page's out-links.
3. **The teleport fix.** A surfer can get stuck on dead-end or cyclic pages. With prob. $1-d$ (≈ 0.15) she **teleports** to a random page. That damping makes the chain *irreducible & aperiodic* — so a *unique* π exists and is reached from any start.
4. **Solve by iterating.** Start uniform, apply M repeatedly — power iteration converges fast on a web-scale, billion-node graph.



A is biggest: three pages point at it. Node size $\propto$ rank.

Why it mattered

Treating links as votes — and ranking by the random walk's stationary distribution — is the idea that launched Google. The same eigenvector trick now drives citation ranking, recommendation graphs, and fraud-ring detection.

THANK YOU

**One coin.
Three scenarios.
A handful of tactics.**

Questions?